

# Data Analysis

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('dataset/final_data.csv')
df.head()
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	hor
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	...	9.0	154
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	...	10.0	102
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0	115

5 rows × 29 columns

	symboling	normalized-losses	make	aspiration	num-of-doors	\		
0	3	122	alfa-romero	std	two			
1	3	122	alfa-romero	std	two			
2	1	122	alfa-romero	std	two			
3	2	164	audi	std	four			
4	2	164	audi	std	four			
	body-style	drive-wheels	engine-location	wheel-base	length	...	\	
0	convertible	rwd	front	88.6	0.811148	...		
1	convertible	rwd	front	88.6	0.811148	...		
2	hatchback	rwd	front	94.5	0.822681	...		
3	sedan	fwd	front	99.8	0.848630	...		
4	sedan	4wd	front	99.4	0.848630	...		
	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg	price	\	
0	9.0	111.0	5000.0	21	27	13495.0		
1	9.0	111.0	5000.0	21	27	16500.0		
2	9.0	154.0	5000.0	19	26	16500.0		
3	10.0	102.0	5500.0	24	30	13950.0		
4	8.0	115.0	5500.0	18	22	17450.0		

## Steps for working with missing data:

1. Identify missing data
2. Deal with missing data
3. Correct data format

```
import numpy as np
```

```
df.replace("?", np.nan, inplace = True)  
df.sample(5)
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression ratio
47	1	104	mazda	std	two	hatchback	fwd	front	93.1	0.764536	...	9.0
165	2	134	toyota	std	two	hatchback	rwd	front	98.4	0.846708	...	9.3
186	3	256	volkswagen	std	two	hatchback	fwd	front	94.5	0.796252	...	8.5
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0
11	0	188	bmw	std	two	sedan	rwd	front	101.2	0.849592	...	9.0

5 rows × 29 columns

```
      symboling  normalized-losses      make aspiration num-of-doors  \  
47           1           104      mazda      std          two  
165          2           134     toyota      std          two  
186          3           256  volkswagen      std          two  
4            2           164       audi      std         four  
11           0           188       bmw      std          two  
  
      body-style drive-wheels engine-location  wheel-base   length   ...  \  
47   hatchback      fwd          front          93.1  0.764536  ...  
165  hatchback      rwd          front          98.4  0.846708  ...  
186  hatchback      fwd          front          94.5  0.796252  ...  
4     sedan        4wd          front          99.4  0.848630  ...  
11     sedan        rwd          front         101.2  0.849592  ...  
  
      compression-ratio  horsepower  peak-rpm  city-mpg  highway-mpg   price  \  
47                9.0         68.0    5000.0      30         31   5195.0  
165               9.3        116.0    4800.0      24         30   9989.0  
186               8.5         90.0    5500.0      24         29   9980.0  
4                 8.0        115.0    5500.0      18         22  17450.0  
11                9.0        121.0    4250.0      21         28  20970.0
```

## Missing Data

The missing values are converted to Python's default. We use Python's built-in functions to identify these missing values. Methods to detect missing data:

1. `.isnull()`
2. `.notnull()`

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
missing_data = df.isnull()
missing_data.head(5)
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	horsepower
0	False	False	False	False	False	False	False	False	False	False	...	False	False
1	False	False	False	False	False	False	False	False	False	False	...	False	False
2	False	False	False	False	False	False	False	False	False	False	...	False	False
3	False	False	False	False	False	False	False	False	False	False	...	False	False
4	False	False	False	False	False	False	False	False	False	False	...	False	False

5 rows × 29 columns

```
      symboling  normalized-losses  make  aspiration  num-of-doors  body-style  \
0          False                False  False        False        False        False
1          False                False  False        False        False        False
2          False                False  False        False        False        False
3          False                False  False        False        False        False
4          False                False  False        False        False        False
```

```
      drive-wheels  engine-location  wheel-base  length  ...  compression-ratio  \
0          False                False        False  False  ...                False
1          False                False        False  False  ...                False
2          False                False        False  False  ...                False
3          False                False        False  False  ...                False
4          False                False        False  False  ...                False
```

```
      horsepower  peak-rpm  city-mpg  highway-mpg  price  city-L/100km  \
0          False    False    False        False  False        False
1          False    False    False        False  False        False
2          False    False    False        False  False        False
3          False    False    False        False  False        False
4          False    False    False        False  False        False
```

"True" stands for missing value, while "False" stands for not missing value.

### Missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values. In the body of the for loop the method ".value\_counts()" counts the number of "True" values.

```
for column in missing_data.columns.values.tolist():
    print(column)
    print(missing_data[column].value_counts())
    print("")
```

symboling

False 201

Name: symboling, dtype: int64

normalized-losses

False 201

Name: normalized-losses, dtype: int64

make

False 201

Name: make, dtype: int64

aspiration

False 201

Name: aspiration, dtype: int64

num-of-doors

False 201

Name: num-of-doors, dtype: int64

Each column has 205 rows of data and 7 columns containing missing data:

1. "normalized-losses": 41 missing data
2. "num-of-doors": 2 missing data
3. "bore": 4 missing data
4. "stroke" : 4 missing data
5. "horsepower": 2 missing data
6. "peak-rpm": 2 missing data

7. "price": 4 missing data

## Deal with missing data

1. drop data
  - a. drop the whole row
  - b. drop the whole column
2. replace data
  - a. replace it by mean
  - b. replace it by frequency
  - c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. We will apply each method to many different columns:

### Replace by mean:

- "normalized-losses": 41 missing data, replace them with mean
- "stroke": 4 missing data, replace them with mean
- "bore": 4 missing data, replace them with mean
- "horsepower": 2 missing data, replace them with mean
- "peak-rpm": 2 missing data, replace them with mean

### Replace by frequency:

- "num-of-doors": 2 missing data, replace them with "four".
  - Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

### Drop the whole row:

- "price": 4 missing data, simply delete the whole row

```
avg_norm = df["normalized-losses"].astype("float").mean(axis=0)
print("Average of normalized-losses:", avg_norm)
```

Average of normalized-losses: 122.0

```
df["normalized-losses"].replace(np.nan, avg_norm, inplace=True)
```

```
avg_of_bore = df["bore"].astype("float").mean(axis=0)  
print("Average of Bore Values:", avg_of_bore)
```

Average of Bore Values: 3.33069156704042

```
df["bore"].replace(np.nan, avg_of_bore, inplace=True)
```

```
avg_stroke = df["stroke"].astype("float").mean(axis=0)  
print("Average of Stroke Values:", avg_of_bore)
```

Average of Stroke Values: 3.33069156704042

```
df["stroke"].replace(np.nan, avg_stroke, inplace = True)
```

```
avg_horsepower = df['horsepower'].astype("float").mean(axis=0)  
print("Average horsepower:", avg_horsepower)
```

Average horsepower: 103.40553390682057

```
df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

```
avg_peak_rpm=df['peak-rpm'].astype('float').mean(axis=0)  
print("Average peak rpm:", avg_peak_rpm)
```

Average peak rpm: 5117.665367742568

```
df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

NameError: name 'avg\_peakrpm' is not defined

To see which values are present in a particular column, use the ".value\_counts()" method:

```
df['num-of-doors'].value_counts()
```

```
four      115
two        86
Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type.

```
df['num-of-doors'].value_counts().idxmax()
```

```
'four'
```

```
#replace the missing 'num-of-doors' values by the most frequent
df['num-of-doors'].replace(np.nan, "four", inplace=True)
```

Finally, let's drop all rows that do not have price data:

```
# simply drop whole row with NaN in "price" column
df.dropna(subset=["price"],axis=0,inplace=True)

# reset index, because we dropped two rows
df.reset_index(drop=True,inplace=True)
```

```
df.head()
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	hor
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	...	9.0	154
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	...	10.0	102
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0	115

5 rows × 29 columns

```

    symboling  normalized-losses      make aspiration num-of-doors \
0           3             122  alfa-romero      std           two
1           3             122  alfa-romero      std           two
2           1             122  alfa-romero      std           two
3           2             164        audi      std          four
4           2             164        audi      std          four

    body-style drive-wheels engine-location  wheel-base  length  ... \
0  convertible      rwd        front        88.6  0.811148  ...
1  convertible      rwd        front        88.6  0.811148  ...
2   hatchback      rwd        front        94.5  0.822681  ...
3        sedan      fwd        front        99.8  0.848630  ...
4        sedan      4wd        front        99.4  0.848630  ...

    compression-ratio  horsepower  peak-rpm  city-mpg  highway-mpg  price \
0                9.0        111.0    5000.0        21          27  13495.0
1                9.0        111.0    5000.0        21          27  16500.0
2                9.0        154.0    5000.0        19          26  16500.0
3               10.0        102.0    5500.0        24          30  13950.0
4                8.0        115.0    5500.0        18          22  17450.0
```



Now, the dataset with no missing values is obtained.

### Convert data types to proper format

```
df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

### Columns after the conversion

df.dtypes

symboling	int64
normalized-losses	int64
make	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64
length	float64
width	float64
height	float64
curb-weight	int64
engine-type	object
num-of-cylinders	object
engine-size	int64
fuel-system	object
bore	float64
stroke	float64
compression-ratio	float64

Finally the cleaned dataset is obtained with no missing values and all data in its proper format.

## Data Standardization

```
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	fuel-system	bore	stroke
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	mpfi	3.47	2.68
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	mpfi	3.47	2.68
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	mpfi	2.68	3.47
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	mpfi	3.19	3.40
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	mpfi	3.19	3.40

5 rows × 27 columns

```
      symboling  normalized-losses      make fuel-type aspiration \
0             3             122  alfa-romero      gas      std
1             3             122  alfa-romero      gas      std
2             1             122  alfa-romero      gas      std
3             2             164        audi      gas      std
4             2             164        audi      gas      std

      num-of-doors  body-style  drive-wheels  engine-location  wheel-base \
0             two  convertible          rwd          front      88.6
1             two  convertible          rwd          front      88.6
2             two   hatchback          rwd          front      94.5
3             four     sedan          fwd          front      99.8
4             four     sedan          4wd          front      99.4

      ...      fuel-system  bore  stroke  compression-ratio  horsepower \
0      ...              mpfi  3.47   2.68              9.0          111
1      ...              mpfi  3.47   2.68              9.0          111
2      ...              mpfi  2.68   3.47              9.0          154
3      ...              mpfi  3.19   3.40             10.0          102
4      ...              mpfi  3.19   3.40              8.0          115
```

```
df["city-L/100km"] = 235 / df["city-mpg"]
```

```
df.head()
```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	hor
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	...	9.0	154
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	...	10.0	102
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0	115

5 rows × 29 columns

```

    symboling  normalized-losses      make aspiration num-of-doors \
0           3             122  alfa-romero      std           two
1           3             122  alfa-romero      std           two
2           1             122  alfa-romero      std           two
3           2             164        audi      std          four
4           2             164        audi      std          four

    body-style drive-wheels engine-location  wheel-base   length   ... \
0  convertible      rwd        front        88.6  0.811148   ...
1  convertible      rwd        front        88.6  0.811148   ...
2   hatchback      rwd        front        94.5  0.822681   ...
3     sedan      fwd        front        99.8  0.848630   ...
4     sedan      4wd        front        99.4  0.848630   ...

    compression-ratio  horsepower  peak-rpm  city-mpg  highway-mpg   price \
0                9.0        111.0    5000.0      21          27  13495.0
1                9.0        111.0    5000.0      21          27  16500.0
2                9.0        154.0    5000.0      19          26  16500.0
3               10.0        102.0    5500.0      24          30  13950.0
4                8.0        115.0    5500.0      18          22  17450.0

```

```
df["highway-mpg"] = 235/df["highway-mpg"]
df.rename(columns={"highway-mpg":'highway-L/100km'}, inplace=True)
```

## Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling variable so the variable values range from 0 to 1

```
# replace (original value) by (original value)/(maximum value)
df["length"] = df["length"]/df["length"].max()
df["width"] = df["width"]/df["width"].max()
df['height'] = df['height']/df['height'].max()
df[["length","width","height"]].head()
```

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

```
      length      width      height
0  0.811148  0.890278  0.816054
1  0.811148  0.890278  0.816054
2  0.822681  0.909722  0.876254
3  0.848630  0.919444  0.908027
4  0.848630  0.922222  0.908027
```

Here we can see, we've normalized "length", "width" and "height" in the range of [0,1].

## Binning

Binning is a process of transforming continuous numerical variables into discrete categorical 'bins', for grouped analysis.

```
df["horsepower"] = df["horsepower"].astype(int,copy=True)
```


Plot the histogram of horsepower in order to see what the distribution of horsepower looks like.

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.hist(df["horsepower"], color='green', edgecolor='black', bins=10)

plt.xlabel("horsepower")
plt.ylabel("count")
plt.title("HORSEPOWER ")
```

Text(0.5, 1.0, 'HORSEPOWER ')

[Download](#)



<Figure size 640x480 with 1 Axes>

```
bins = np.linspace(min(df["horsepower"]),max(df["horsepower"]),4)
bins
```

```
array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names ,include_lowes
df[['horsepower', 'horsepower-binned']].head(10)
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low

```
horsepower horsepower-binned
0          111              Low
1          111              Low
2          154            Medium
3          102              Low
4          115              Low
```

5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low

Lets see the number of vehicles in each bin.

```
df["horsepower-binned"].value_counts()
```

```
Low      153
Medium   43
High      5
Name: horsepower-binned, dtype: int64
```


Lets plot the distribution of each bin.

```
%matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.barh(group_names, df["horsepower-binned"].value_counts() , color='red', edgecolor='

plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
```

```
Text(0, 0.5, 'count')
```

[Download](#)



<Figure size 640x480 with 1 Axes>

```
df.corr()
```



	symboling	normalized-losses	wheel-base	length	width	height	curb-weight	engine-size	bore	stroke
symboling	1.000000	0.466264	-0.535987	-0.365404	-0.242423	-0.550160	-0.233118	-0.110581	-0.140019	-0.007992
normalized-losses	0.466264	1.000000	-0.056661	0.019424	0.086802	-0.373737	0.099404	0.112360	-0.029862	0.055127
wheel-base	-0.535987	-0.056661	1.000000	0.876024	0.814507	0.590742	0.782097	0.572027	0.493244	0.157964
length	-0.365404	0.019424	0.876024	1.000000	0.857170	0.492063	0.880665	0.685025	0.608971	0.123913
width	-0.242423	0.086802	0.814507	0.857170	1.000000	0.306002	0.866201	0.729436	0.544885	0.188814
height	-0.550160	-0.373737	0.590742	0.492063	0.306002	1.000000	0.307581	0.074694	0.180449	-0.060822
curb-weight	-0.233118	0.099404	0.782097	0.880665	0.866201	0.307581	1.000000	0.849072	0.644060	0.167412
engine-size	-0.110581	0.112360	0.572027	0.685025	0.729436	0.074694	0.849072	1.000000	0.572609	0.205806
bore	-0.140019	-0.029862	0.493244	0.608971	0.544885	0.180449	0.644060	0.572609	1.000000	-0.055390
stroke	-0.007992	0.055127	0.157964	0.123913	0.188814	-0.060822	0.167412	0.205806	-0.055390	1.000000
compression-ratio	-0.182196	-0.114713	0.250313	0.159733	0.189867	0.259737	0.156433	0.028889	0.001263	0.187854
horsepower	0.075819	0.217299	0.371147	0.579821	0.615077	-0.087027	0.757976	0.822676	0.566936	0.098282
peak-rpm	0.279740	0.239543	-0.360305	-0.285970	-0.245800	-0.309974	-0.279361	-0.256733	-0.267392	-0.063388
city-mpg	-0.035527	-0.225016	-0.470606	-0.665192	-0.633531	-0.049800	-0.749543	-0.650546	-0.582027	-0.034079
highway-mpg	0.036233	-0.181877	-0.543304	-0.698142	-0.680635	-0.104812	-0.794889	-0.679571	-0.591309	-0.034741
price	-0.082391	0.133999	0.584642	0.690628	0.751265	0.135486	0.834415	0.872335	0.543155	0.082267
city-L/100km	0.066171	0.238567	0.476153	0.657373	0.673363	0.003811	0.785353	0.745059	0.554610	0.036285
diesel	-0.196735	-0.101546	0.307237	0.211187	0.244356	0.281578	0.221046	0.070779	0.054458	0.241033
gas	0.196735	0.101546	-0.307237	-0.211187	-0.244356	-0.281578	-0.221046	-0.070779	-0.054458	-0.241033

	symboling	normalized-losses	wheel-base	length	\
symboling	1.000000	0.466264	-0.535987	-0.365404	
normalized-losses	0.466264	1.000000	-0.056661	0.019424	
wheel-base	-0.535987	-0.056661	1.000000	0.876024	
length	-0.365404	0.019424	0.876024	1.000000	
width	-0.242423	0.086802	0.814507	0.857170	
height	-0.550160	-0.373737	0.590742	0.492063	
curb-weight	-0.233118	0.099404	0.782097	0.880665	
engine-size	-0.110581	0.112360	0.572027	0.685025	
bore	-0.140019	-0.029862	0.493244	0.608971	
stroke	-0.007992	0.055127	0.157964	0.123913	
compression-ratio	-0.182196	-0.114713	0.250313	0.159733	
horsepower	0.075819	0.217299	0.371147	0.579821	
peak-rpm	0.279740	0.239543	-0.360305	-0.285970	
city-mpg	-0.035527	-0.225016	-0.470606	-0.665192	
highway-mpg	0.036233	-0.181877	-0.543304	-0.698142	
price	-0.082391	0.133999	0.584642	0.690628	
city-L/100km	0.066171	0.238567	0.476153	0.657373	
diesel	-0.196735	-0.101546	0.307237	0.211187	
gas	0.196735	0.101546	-0.307237	-0.211187	

```
/tmp/ipykernel_6935/1134722465.py:1: FutureWarning: The default value of numeric_only in
df.corr()
```

```
df[['bore', 'stroke', 'compression-ratio', 'horsepower']].corr()
```

	bore	stroke	compression-ratio	horsepower
bore	1.000000	-0.055390	0.001263	0.566936
stroke	-0.055390	1.000000	0.187854	0.098282
compression-ratio	0.001263	0.187854	1.000000	-0.214514
horsepower	0.566936	0.098282	-0.214514	1.000000

	bore	stroke	compression-ratio	horsepower
bore	1.000000	-0.055390	0.001263	0.566936
stroke	-0.055390	1.000000	0.187854	0.098282
compression-ratio	0.001263	0.187854	1.000000	-0.214514
horsepower	0.566936	0.098282	-0.214514	1.000000


## Positive linear relationship

Scatterplot of "engine-size" and "price"

```
sns.regplot(x="engine-size", y="price", data=df ,color='green')
plt.ylim(0,)
```

(0.0, 53417.19589840475)

[Download](#)



<Figure size 640x480 with 1 Axes>

```
df[["engine-size", "price"]].corr()
```

	engine-size	price
engine-size	1.000000	0.872335
price	0.872335	1.000000


```

engine-size    engine-size    price
engine-size    1.000000    0.872335
price          0.872335    1.000000

```

Thus, Highway mpg is a potential predictor variable of price

```
sns.regplot(x="highway-mpg", y="price", data=df, color='grey');
```



<Figure size 640x480 with 1 Axes>

```
df[['highway-mpg', 'price']].corr()
```

	highway-mpg	price
highway-mpg	1.000000	-0.704692
price	-0.704692	1.000000

```


          highway-mpg    price
highway-mpg    1.000000 -0.704692
price          -0.704692  1.000000

```

## Weak Linear Relationship

Let's see if "Peak-rpm" as a predictor variable of "price".

```
sns.regplot(x="peak-rpm", y="price", data=df ,color='blue');
```




<Figure size 640x480 with 1 Axes>

## Categorical variables

Relationship between "body-style" and "price"

```
sns.boxplot(x="body-style", y="price", data=df);
```


[Download](#)



<Figure size 640x480 with 1 Axes>

```
sns.boxplot(x="engine-location", y="price", data=df);
```

[Download](#)



<Figure size 640x480 with 1 Axes>

```
df.describe(include=['object'])
```

	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	engine-type	num-of-cylinders	fuel-system	horsepower-binned
count	201	201	201	201	201	201	201	201	201	200
unique	22	2	2	5	3	2	6	7	8	3
top	toyota	std	four	sedan	fwd	front	ohc	four	mpfi	Low
freq	32	165	115	94	118	198	145	157	92	115

```

count      make aspiration num-of-doors body-style drive-wheels \
unique      22         2         2         5         3
top      toyota      std      four      sedan      fwd
freq       32      165      115      94      118

```

```

count      engine-location engine-type num-of-cylinders fuel-system \
unique         2         6         7         8
top          front      ohc      four      mpfi
freq         198      145      157      92

```

```

count      horsepower-binned
unique         3

```

top	Low
freq	115

### P-value:

What is this P-value? The P-value is the probability value that the correlation between these two variables is statistically significant. Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the

- p-value is  $< 0.001$ : we say there is strong evidence that the correlation is significant.
- the p-value is  $< 0.05$ : there is moderate evidence that the correlation is significant.
- the p-value is  $< 0.1$ : there is weak evidence that the correlation is significant.
- the p-value is  $> 0.1$ : there is no evidence that the correlation is significant.

```
from scipy import stats
```

## Wheel-base vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'wheel-base' and 'price'.

```
pearson_coef, p_value = stats.pearsonr(df['wheel-base'], df['price'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_
```

The Pearson Correlation Coefficient is 0.584641822265508 with a P-value of P = 8.076488

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between wheel-base and price is statistically significant, although the linear relationship isn't extremely strong ( $\sim 0.585$ )



## Horse-Power vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'Horse-Power' and 'price'.

```
pearson_coef, p_value = stats.pearsonr(df['horsepower'], df['price'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_
```

---

The Pearson Correlation Coefficient is 0.8095745670036562 with a P-value of P = 6.36905

---

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between horsepower and price is statistically significant, and the linear relationship is quite strong (~0.809, close to 1)

## Length vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'length' and 'price'.

```
pearson_coef, p_value = stats.pearsonr(df['length'], df['price'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p
```

---

The Pearson Correlation Coefficient is 0.690628380448364 with a P-value of P = 8.01647

---

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between length and price is statistically significant, and the linear relationship is moderately strong ( $\sim 0.691$ ).

## Width vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'width' and 'price':

```
pearson_coef, p_value = stats.pearsonr(df['width'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_
```

```
The Pearson Correlation Coefficient is 0.7512653440522674  with a P-value of P = 9.20033
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between width and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.751$ ).

## Curb-weight vs Price

```
pearson_coef, p_value = stats.pearsonr(df['curb-weight'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =",
```

```
The Pearson Correlation Coefficient is 0.8344145257702846  with a P-value of P = 2.1895
```

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between curb-weight and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.834$ ).

## Engine-size vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'engine-size' and 'price':

```
pearson_coef, p_value = stats.pearsonr(df['engine-size'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_
```

The Pearson Correlation Coefficient is 0.8723351674455185 with a P-value of P = 9.26549

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between engine-size and price is statistically significant, and the linear relationship is very strong ( $\sim 0.872$ ).

## Bore vs Price

```
pearson_coef, p_value = stats.pearsonr(df['bore'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ",
```

The Pearson Correlation Coefficient is 0.5431553832626602 with a P-value of P = 8.049

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between bore and price is statistically significant, but the linear relationship is only moderate ( $\sim 0.521$ ).

## City-mpg vs Price

```
pearson_coef, p_value = stats.pearsonr(df['city-mpg'], df['price'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p
```

The Pearson Correlation Coefficient is -0.6865710067844677 with a P-value of P = 2.321

### Conclusion:

Since the p-value is  $< 0.001$ , the correlation between city-mpg and price is statistically significant, and the coefficient of  $\sim -0.687$  shows that the relationship is negative and moderately strong.

## Conclusion: Important Variables

We now have a better idea of what our data looks like and which variables are important to take into account when predicting the car price. We have narrowed it down to the following variables:

Continuous numerical variables:

- Length
- Width
- Curb-weight
- Engine-size
- Horsepower
- City-mpg
- Highway-mpg
- Wheel-base

- Bore

Categorical variables:

- Drive-wheels

As we now move into building machine learning models to automate our analysis, feeding the model with variables that meaningfully affect our target variable will improve our model's prediction performance.