

```
import pandas as pd
```

```
df = pd.read_csv('dataset/final_data.csv')
```

```
df.head(5)
```

```

      symboling  normalized-losses      make aspiration num-of-doors \
0           3           122  alfa-romero      std           two
1           3           122  alfa-romero      std           two
2           1           122  alfa-romero      std           two
3           2           164      audi      std           four
4           2           164      audi      std           four

      body-style drive-wheels engine-location  wheel-base  length  ... \
0  convertible      rwd      front      88.6  0.811148  ...
1  convertible      rwd      front      88.6  0.811148  ...
2   hatchback      rwd      front      94.5  0.822681  ...
3      sedan      fwd      front      99.8  0.848630  ...
4      sedan      4wd      front      99.4  0.848630  ...

      compression-ratio  horsepower  peak-rpm  city-mpg  highway-mpg  price \
0           9.0      111.0    5000.0      21      27  13495.0
1           9.0      111.0    5000.0      21      27  16500.0
2           9.0      154.0    5000.0      19      26  16500.0
3          10.0      102.0    5500.0      24      30  13950.0
4           8.0      115.0    5500.0      18      22  17450.0

```

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	hor
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	...	9.0	154
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	...	10.0	102
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0	115

5 rows × 29 columns

```
df.columns
```

```
Index(['symboling', 'normalized-losses', 'make', 'aspiration', 'num-of-doors',  
      'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',  
      'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',  
      'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',  
      'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price',  
      'city-L/100km', 'horsepower-binned', 'diesel', 'gas'],  
      dtype='object')
```

```
X = df[['length', 'width', 'curb-weight', 'engine-size', 'horsepower', 'city-mpg', 'highway-mpg']  
Y = df[['price']].copy()
```

X

	length	width	curb-weight	engine-size	horsepower	city-mpg	\
0	0.811148	0.890278	2548	130	111.0	21	
1	0.811148	0.890278	2548	130	111.0	21	
2	0.822681	0.909722	2823	152	154.0	19	
3	0.848630	0.919444	2337	109	102.0	24	
4	0.848630	0.922222	2824	136	115.0	18	
..	
196	0.907256	0.956944	2952	141	114.0	23	
197	0.907256	0.955556	3049	141	160.0	19	
198	0.907256	0.956944	3012	173	134.0	18	
199	0.907256	0.956944	3217	145	106.0	26	
200	0.907256	0.956944	3062	141	114.0	19	

	highway-mpg	wheel-base	bore
0	27	88.6	3.47
1	27	88.6	3.47
2	26	94.5	2.68
3	30	99.8	3.19
4	22	99.4	3.19
..

	length	width	curb-weight	engine-size	horsepower	city-mpg	highway-mpg	wheel-base	bore
0	0.811148	0.890278	2548	130	111.0	21	27	88.6	3.47
1	0.811148	0.890278	2548	130	111.0	21	27	88.6	3.47
2	0.822681	0.909722	2823	152	154.0	19	26	94.5	2.68
3	0.848630	0.919444	2337	109	102.0	24	30	99.8	3.19
4	0.848630	0.922222	2824	136	115.0	18	22	99.4	3.19
...
196	0.907256	0.956944	2952	141	114.0	23	28	109.1	3.78
197	0.907256	0.955556	3049	141	160.0	19	25	109.1	3.78
198	0.907256	0.956944	3012	173	134.0	18	23	109.1	3.58
199	0.907256	0.956944	3217	145	106.0	26	27	109.1	3.01
200	0.907256	0.956944	3062	141	114.0	19	25	109.1	3.78

201 rows × 9 columns

Y

```

      price
0    13495.0
1    16500.0
2    16500.0
3    13950.0
4    17450.0
..      ...
196  16845.0
197  19045.0
198  21485.0
199  22470.0
200  22625.0

```

[201 rows x 1 columns]

	price
0	13495.0
1	16500.0
2	16500.0
3	13950.0
4	17450.0
...	...
196	16845.0
197	19045.0
198	21485.0
199	22470.0
200	22625.0

201 rows × 1 columns

```
feature = X.values
```

```
value = Y.values
```

```
from sklearn.model_selection import train_test_split
feature_train, feature_test, value_train, value_test = train_test_split(feature, value, test_size
```

Feature Scaling

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
feature_train = sc.fit_transform(feature_train)
feature_test = sc.fit_transform(feature_test)
```

```
from sklearn.linear_model import LinearRegression
# create an object
ml=LinearRegression()
ml.fit(feature_train,value_train)
```

LinearRegression()

▼ LinearRegression
LinearRegression()

```
value_pred = ml.predict(feature_test)
print(value_pred)
```

```
[[ 5844.78826709]
 [11348.76667238]
 [20251.35191747]
 [21406.74779824]
 [20982.97703476]
 [ 9879.57644106]
 [14505.61640274]
 [ 6268.36622588]
 [17371.50852175]
 [ 7133.70247958]
 [11507.71500216]
 [20541.38134988]
 [ 7886.8953474 ]
 [ 8265.97639025]
 [17609.8386963 ]
 [18850.71526732]
 [ 7187.296721 ]
 [15162.619681 ]
 [10059.0243429 ]
 [ 6519.28263493]
```

```
value=[[0.811148,0.890278,2548,130,111.0,21,27,88.6,3.47]]
```

```
ml.predict(value)
```

```
array([[11136.98548867]])
```

```
from sklearn.metrics import r2_score
r2_score(value_test,value_pred)
```


0.7526252480045523

```
import matplotlib.pyplot as plt
plt.figure(figsize=(12,10))
plt.scatter(value_test,value_pred)
plt.xlabel('Actual')
plt.ylabel('Predict')
```

Text(0, 0.5, 'Predict')

<Figure size 864x720 with 1 Axes>

[Download](#)



```
future_value = pd.DataFrame({'Actual value':value_test, 'Predicted value':value_pred, 'Difference': value_test-  
value_pred}) future_value[0:20]
```

```
import pickle  
pickle.dump(LinearRegression,open("model.pkl", "wb"))
```